

# CapstoneProject1

August 16, 2020

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.preprocessing import StandardScaler, normalize
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
[13]: df_db1 = pd.read_csv('health care diabetes.csv')
```

```
[14]: df_db1.head(10)
```

```
[14]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
5	5	116	74	0	0	25.6	
6	3	78	50	32	88	31.0	
7	10	115	0	0	0	35.3	
8	2	197	70	45	543	30.5	
9	8	125	96	0	0	0.0	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0
6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1

```
[15]: df_db1.shape
```

```
[15]: (768, 9)
```

```
[16]: df_db1.dtypes
```

```
[16]: Pregnancies          int64
      Glucose             int64
      BloodPressure       int64
      SkinThickness       int64
      Insulin             int64
      BMI                 float64
      DiabetesPedigreeFunction float64
      Age                 int64
      Outcome             int64
      dtype: object
```

```
[17]: df_db1.describe()
```

```
[17]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[18]: df_db1.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
```

1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)  
memory usage: 54.1 KB


```
[19]: sns.distplot(df_db1['Pregnancies'])
```

```
[19]: <AxesSubplot:xlabel='Pregnancies'>
```




```
[20]: sns.distplot(df_db1['Glucose'])
```

```
[20]: <AxesSubplot:xlabel='Glucose'>
```




```
[21]: sns.distplot(df_db1['BloodPressure'])
```

```
[21]: <AxesSubplot:xlabel='BloodPressure'>
```




```
[22]: sns.distplot(df_db1['SkinThickness'])
```

```
[22]: <AxesSubplot:xlabel='SkinThickness'>
```




```
[23]: sns.distplot(df_db1['Insulin'])
```

```
[23]: <AxesSubplot:xlabel='Insulin'>
```




```
[24]: sns.distplot(df_db1['BMI'])
```

```
[24]: <AxesSubplot:xlabel='BMI'>
```




```
[25]: sns.distplot(df_db1['DiabetesPedigreeFunction'])
```

```
[25]: <AxesSubplot:xlabel='DiabetesPedigreeFunction'>
```



```
[26]: sns.distplot(df_db1['Age'])
```

```
[26]: <AxesSubplot:xlabel='Age'>
```



```
[28]: print("Skewness: %f" % df_db1['Pregnancies'].skew())
      print("Kurtosis: %f" % df_db1['Pregnancies'].kurt())
```

```
Skewness: 0.901674
Kurtosis: 0.159220
```

```
[29]: print("Skewness: %f" % df_db1['Glucose'].skew())
      print("Kurtosis: %f" % df_db1['Glucose'].kurt())
```

```
Skewness: 0.173754
Kurtosis: 0.640780
```

```
[30]: print("Skewness: %f" % df_db1['BloodPressure'].skew())
      print("Kurtosis: %f" % df_db1['BloodPressure'].kurt())
```

```
Skewness: -1.843608
Kurtosis: 5.180157
```

```
[31]: print("Skewness: %f" % df_db1['SkinThickness'].skew())
      print("Kurtosis: %f" % df_db1['SkinThickness'].kurt())
```

```
Skewness: 0.109372
Kurtosis: -0.520072
```



```
[32]: print("Skewness: %f" % df_db1['Insulin'].skew())
      print("Kurtosis: %f" % df_db1['Insulin'].kurt())
```

```
Skewness: 2.272251
Kurtosis: 7.214260
```

```
[33]: print("Skewness: %f" % df_db1['BMI'].skew())
      print("Kurtosis: %f" % df_db1['BMI'].kurt())
```

```
Skewness: -0.428982
Kurtosis: 3.290443
```

```
[34]: print("Skewness: %f" % df_db1['DiabetesPedigreeFunction'].skew())
      print("Kurtosis: %f" % df_db1['DiabetesPedigreeFunction'].kurt())
```

```
Skewness: 1.919911
Kurtosis: 5.594954
```

```
[35]: print("Skewness: %f" % df_db1['Age'].skew())
      print("Kurtosis: %f" % df_db1['Age'].kurt())
```


```
Skewness: 1.129597
Kurtosis: 0.643159
```

```
[37]: # First, let's count the number of null values
      total = df_db.isnull().sum().sort_values(ascending=False)
      # Then, let's calculate the percentage of missing data per feature
      percent = (df_db1.isnull().sum()/df_db.isnull().count()).
        ↪sort_values(ascending=False)
      # Finally, let's concatenate Total and Percent into another dataframe
      missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
      missing_data.head(20)
```

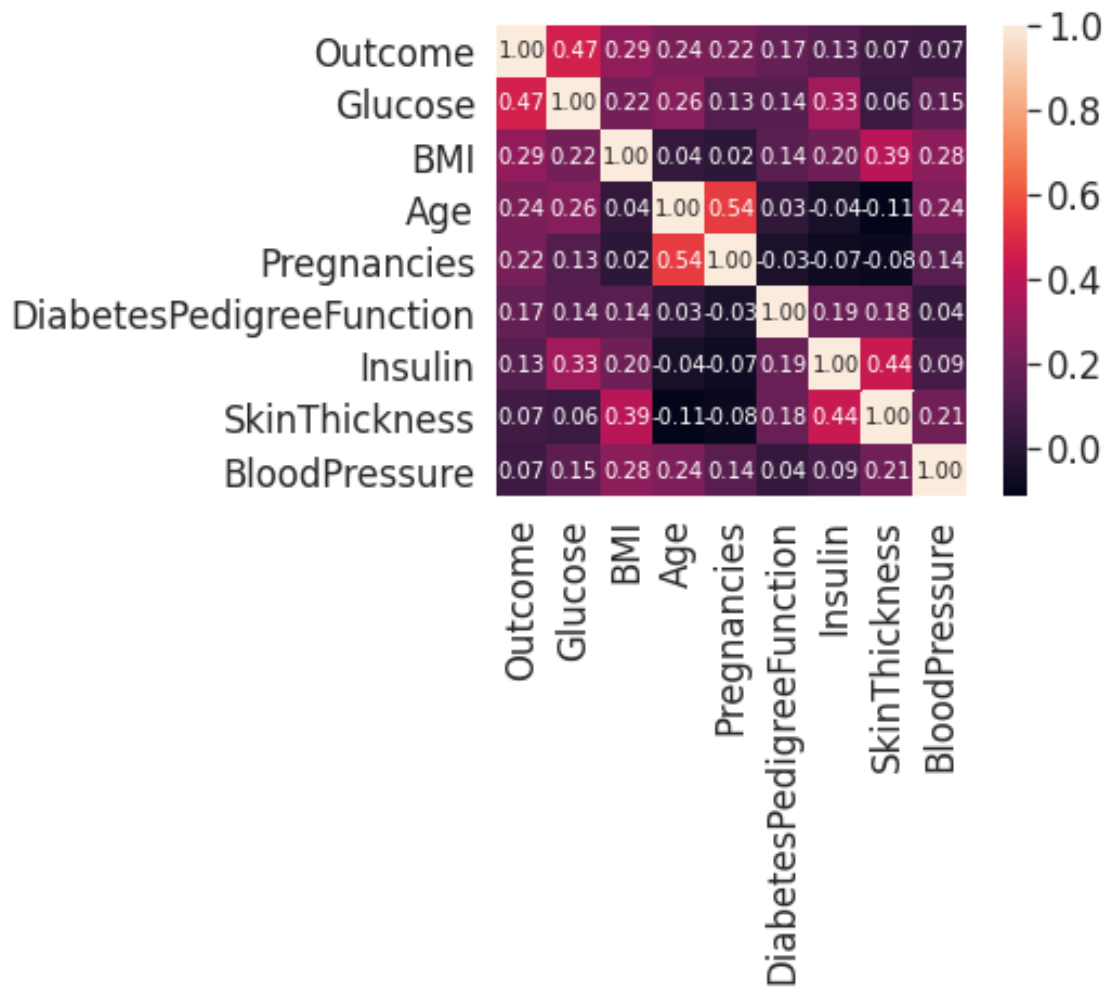
```
[37]:
```

	Total	Percent
Outcome	0	0.0
Age	0	0.0
DiabetesPedigreeFunction	0	0.0
BMI	0	0.0
Insulin	0	0.0
SkinThickness	0	0.0
BloodPressure	0	0.0
Glucose	0	0.0
Pregnancies	0	0.0


```
[38]: corrmat = df_db.corr()
      f, ax = plt.subplots(figsize=(12, 12))
      sns.heatmap(corrmat, vmax=.8, square=True);
```



```
[39]: k = 10 #number of variables for heatmap
cols = corrmat.nlargest(k, 'Outcome')['Outcome'].index
cm = np.corrcoef(df_db1[cols].values.T)
sns.set(font_scale=1.5)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f',
    ↳annot_kws={'size': 10}, yticklabels=cols.values, xticklabels=cols.values)
plt.show()
```




```
[42]: #pairplot
sns.set()
cols = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']
sns.pairplot(df_db[cols], size = 2.5)
plt.show();
```




```
[43]: sns.countplot(df_db1['Pregnancies'])
```

```
[43]: <AxesSubplot:xlabel='Pregnancies', ylabel='count'>
```



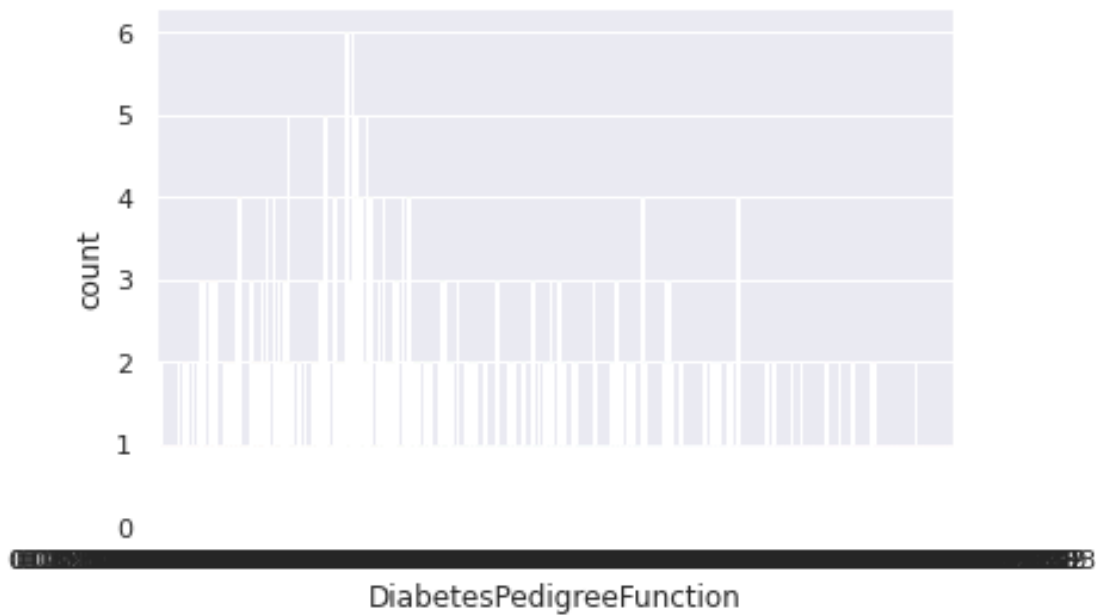
```
[45]: sns.countplot(df_db1['Glucose'])
```

```
[45]: <AxesSubplot:xlabel='Glucose', ylabel='count'>
```




```
[44]: sns.countplot(df_db1['DiabetesPedigreeFunction'])
```

```
[44]: <AxesSubplot:xlabel='DiabetesPedigreeFunction', ylabel='count'>
```



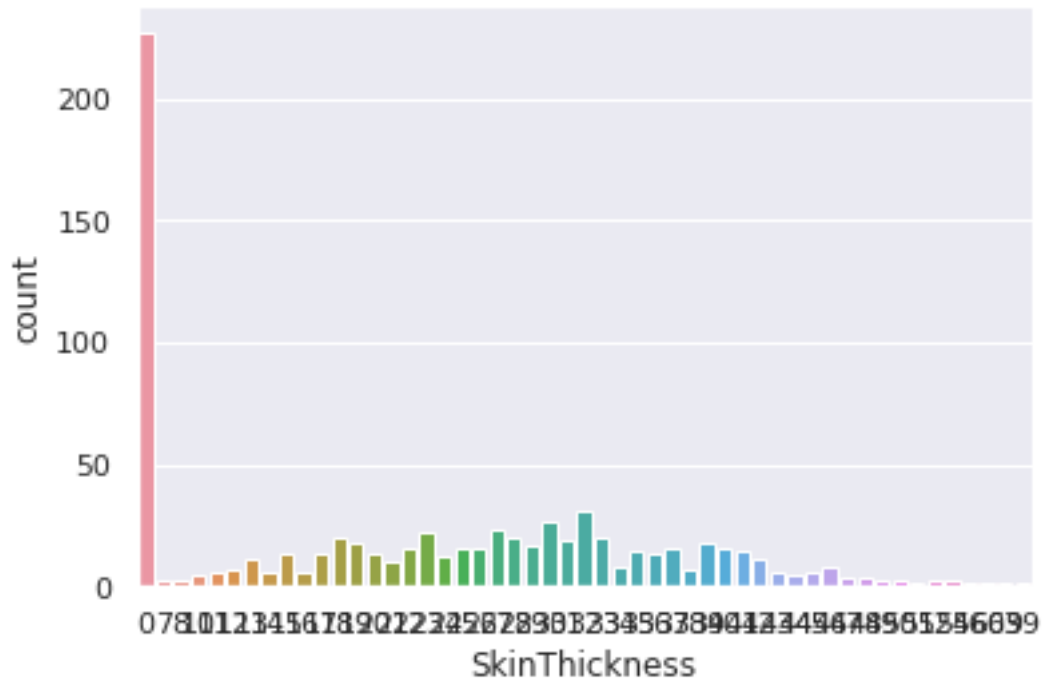
```
[46]: sns.countplot(df_db1['BloodPressure'])
```

```
[46]: <AxesSubplot:xlabel='BloodPressure', ylabel='count'>
```




```
[47]: sns.countplot(df_db1['SkinThickness'])
```

```
[47]: <AxesSubplot:xlabel='SkinThickness', ylabel='count'>
```



```
[67]: sns.countplot(df_db1['BMI'])
```


```
[67]: <AxesSubplot:xlabel='BMI', ylabel='count'>
```



```
[68]: sns.countplot(df_db1['Age'])
```

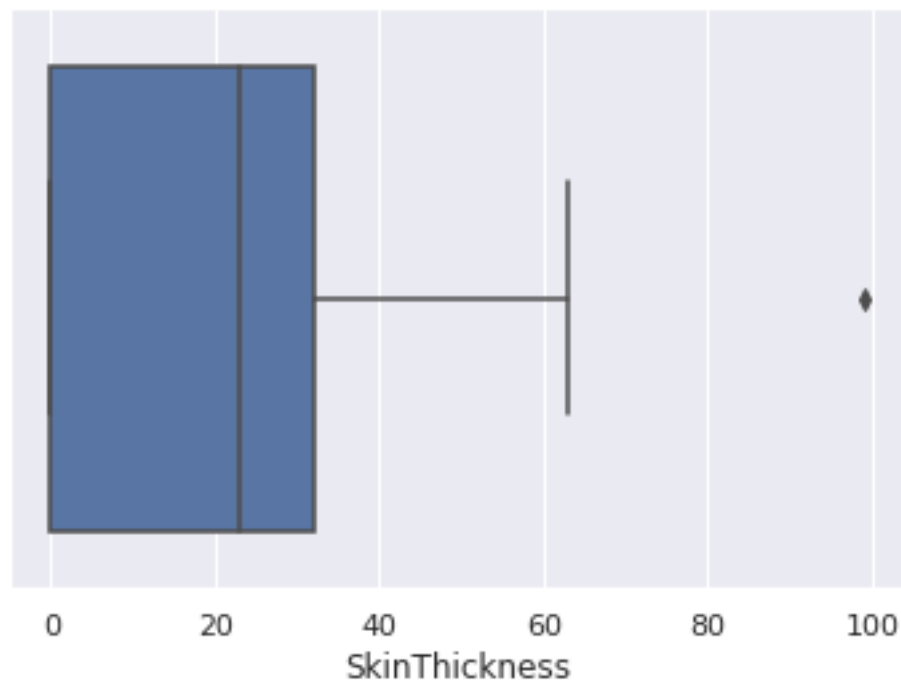
```
[68]: <AxesSubplot:xlabel='Age', ylabel='count'>
```






```
[69]: sns.boxplot(df_db1['SkinThickness'])
```

```
[69]: <AxesSubplot:xlabel='SkinThickness'>
```




```
[70]: sns.boxplot(df_db1['Age'])
```

```
[70]: <AxesSubplot:xlabel='Age'>
```




```
[71]: sns.boxplot(df_db1['BMI'])
```

```
[71]: <AxesSubplot:xlabel='BMI'>
```



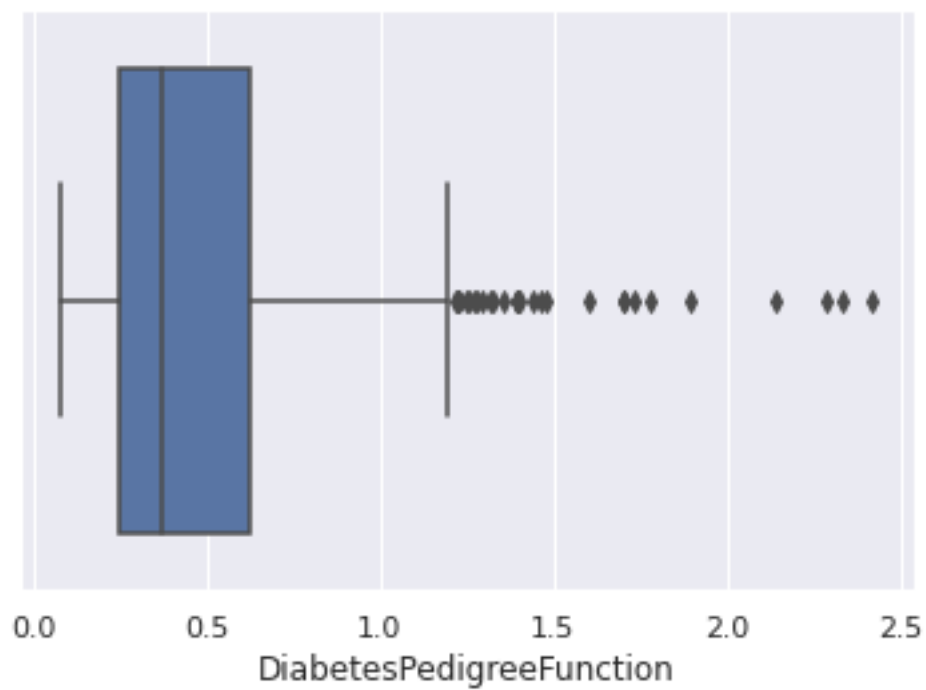
```
[72]: sns.boxplot(df_db1['SkinThickness'])
```

```
[72]: <AxesSubplot:xlabel='SkinThickness'>
```




```
[74]: sns.boxplot(df_db1['DiabetesPedigreeFunction'])
```

```
[74]: <AxesSubplot:xlabel='DiabetesPedigreeFunction'>
```




```
[75]: sns.boxplot(df_db1['Glucose'])
```

```
[75]: <AxesSubplot:xlabel='Glucose'>
```



```
[76]: sns.boxplot(df_db1['Pregnancies'])
```

```
[76]: <AxesSubplot:xlabel='Pregnancies'>
```



```
[79]: from sklearn.model_selection import train_test_split
X = df_db.drop(columns = 'Outcome')
y = df_db['Outcome']
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.
↪2,random_state = 42)
```

```
[92]: print(X_train)
```

```
[[ -0.52639686 -1.15139792 -3.75268255 ... -4.13525578 -0.49073479
  -1.03594038]
 [ 1.58804586 -0.27664283  0.68034485 ... -0.48916881  2.41502991
  1.48710085]
 [-0.82846011  0.56687102 -1.2658623  ... -0.42452187  0.54916055
 -0.94893896]
 ...
 [ 1.8901091  -0.62029661  0.89659009 ...  1.76054443  1.981245
  0.44308379]
 [-1.13052335  0.62935353 -3.75268255 ...  1.34680407 -0.78487662
 -0.33992901]
 [-1.13052335  0.12949347  1.43720319 ... -1.22614383 -0.61552223
 -1.03594038]]
```

```
[89]: from sklearn.preprocessing import StandardScaler
scalar = StandardScaler()
train_t = scalar.fit_transform (X_train) #fit and transform
test_t = scalar.transform (X_test) # only transform
print(train_t)
```

```
[[ -0.52639686 -1.15139792 -3.75268255 ... -4.13525578 -0.49073479
  -1.03594038]
 [ 1.58804586 -0.27664283  0.68034485 ... -0.48916881  2.41502991
  1.48710085]
 [-0.82846011  0.56687102 -1.2658623  ... -0.42452187  0.54916055
 -0.94893896]
 ...
 [ 1.8901091  -0.62029661  0.89659009 ...  1.76054443  1.981245
  0.44308379]
 [-1.13052335  0.62935353 -3.75268255 ...  1.34680407 -0.78487662
 -0.33992901]
 [-1.13052335  0.12949347  1.43720319 ... -1.22614383 -0.61552223
 -1.03594038]]
```

```
[93]: #Fitting the Model to the training data
from sklearn.dummy import DummyClassifier
from sklearn.linear_model import LogisticRegression
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
import xgboost

```

```

[94]: from sklearn.model_selection import KFold, cross_val_score
for model in [
    DummyClassifier,
    LogisticRegression,
    DecisionTreeClassifier,
    KNeighborsClassifier,
    GaussianNB,
    SVC,
    RandomForestClassifier,
    xgboost.XGBClassifier,
]:
    cls = model()
    kf = KFold(n_splits = 5, random_state = 45)
    score = cross_val_score(cls, train_t, y_train, cv = kf, scoring = "roc_auc")
    print(f" {model.__name__:22} AUC:"
          f"\t {score.mean():.3f} STD: {score.std():.2f}")

```

DummyClassifier	AUC:	0.481	STD: 0.02
LogisticRegression	AUC:	0.832	STD: 0.02
DecisionTreeClassifier	AUC:	0.674	STD: 0.03
KNeighborsClassifier	AUC:	0.786	STD: 0.04
GaussianNB	AUC:	0.803	STD: 0.04
SVC	AUC:	0.833	STD: 0.04
RandomForestClassifier	AUC:	0.835	STD: 0.03
XGBClassifier	AUC:	0.809	STD: 0.03

```

[95]: #Without any Hyper parameters we see "RandomForestClassifier" model has the
      ↪ best accuracy
      #making the model work
      rfe = RandomForestClassifier(n_estimators = 1000, random_state = 42)

```

```

[96]: rfe.fit(train_t, y_train)

```

```

[96]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=None, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=1000,
                             n_jobs=None, oob_score=False, random_state=42, verbose=0,

```

```
warm_start=False)
```

```
[97]: #Evaluating the model
print("Accuracy on test set is", round(rfe.score(test_t, y_test)*100,2),"%")
```

Accuracy on test set is 74.03 %

```
[98]: from sklearn.metrics import precision_score
print("Precision score is", round(precision_score(y_test, rfe.
    ↳predict(test_t))*100,2))
```

Precision score is 63.16

```
[105]: print("feature importance is \n")
for i,j in zip(X.columns.tolist(), rfe.feature_importances_.tolist()):
    if(i=="DiabetesPedigreeFunction"):
        print(i, "\t\t",j)
    elif(i=="Age" or i=="BMI"):
        print(i, "\t\t\t\t\t", j)
    else:
        print(i, "\t\t\t\t", j)
```

feature importance is

Pregnancies	0.07870476973503568
Glucose	0.257006690994012
BloodPressure	0.08838587879978824
SkinThickness	0.0682424578137327
Insulin	0.07746893645810259
BMI	0.16419653560403633
DiabetesPedigreeFunction	0.12046534553435637
Age	0.14552938506093616

```
[106]: #lets do hyperparam tuning
from sklearn.model_selection import GridSearchCV
new_rfe = RandomForestClassifier()
params = {
    "max_features": [0.4, "auto"],
    "n_estimators": [15, 200, 500, 1000],
    "min_samples_leaf": [1, 0.1],
    "random_state": [42],
}
cvs = GridSearchCV(new_rfe, params, n_jobs = -1).fit(train_t, y_train)
```

```
[107]: print(cvs.best_score_)
```

0.7785285885645742



```
[108]: print(cvs.best_estimator_)
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=200,
                        n_jobs=None, oob_score=False, random_state=42, verbose=0,
                        warm_start=False)
```

```
[109]: print(cvs.best_params_)
```

```
{'max_features': 'auto', 'min_samples_leaf': 1, 'n_estimators': 200,
 'random_state': 42}
```

```
[115]: #fitting the model with best params
rfe_final = RandomForestClassifier()
params = {
    "max_features":["auto"],
    "n_estimators":[200],
    "min_samples_leaf":[1],
    "random_state":[42],
}
```

```
[116]: rfe_final.fit(train_t, y_train)
```

```
[116]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                              criterion='gini', max_depth=None, max_features='auto',
                              max_leaf_nodes=None, max_samples=None,
                              min_impurity_decrease=0.0, min_impurity_split=None,
                              min_samples_leaf=1, min_samples_split=2,
                              min_weight_fraction_leaf=0.0, n_estimators=100,
                              n_jobs=None, oob_score=False, random_state=None,
                              verbose=0, warm_start=False)
```

```
[117]: y_pred = rfe_final.predict(test_t)
```

```
[118]: print(precision_score(y_test, y_pred))
```

```
0.6363636363636364
```

```
[120]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.80	0.80	99

1	0.64	0.64	0.64	55
accuracy			0.74	154
macro avg	0.72	0.72	0.72	154
weighted avg	0.74	0.74	0.74	154

```
[121]: #print Roc_auc_score
from sklearn.metrics import roc_auc_score
print(roc_auc_score(y_test, y_pred))
```

0.7171717171717171

```
[ ]:
```